

From Narrative Inflation to Verifiable Self-Correction: An Empirical Longitudinal Case Study of Long-Horizon LLM Mathematical Reasoning on the Riemann Hypothesis

從敘事膨脹到可驗證自我修正：大語言模型在黎曼猜想長程多輪數學推理中的失敗模式與修正機制實證案例研究

Riemann Hypothesis Research Collective (AI-Human Collaboration)

Preprint • Project Scope: Rounds 1–388 (145 Formal Peer-Review Audits) • August 2026


Abstract—Current benchmarks for evaluating mathematical reasoning in Large Language Models (LLMs) predominantly focus on static, single-turn, closed-form contest problems with known ground truths. Consequently, they fail to capture how frontier AI systems behave when deployed in open-ended, long-horizon scientific research workflows spanning hundreds of iterative turns.

In this paper, we present an extensive empirical longitudinal case study spanning **388 research rounds and 145 formal adversarial peer-review cycles** over a human-AI collaborative effort attacking the Riemann Hypothesis (RH) via symplectic Dirac operator theory. We discover, analyze, and formalize a **10-class taxonomy of multi-turn mathematical reasoning failure modes**, ranging from coordinate confusion, hidden circularity, and ensemble/pointwise category mixing to unchecked perturbation expansion domains and narrative progress inflation.

Crucially, we uncover the exact dynamical mechanism of **Multi-Turn Epistemic Drift and Adversarial Self-Correction**: while vague prompts invariably induce models to escalate narrative bluffing, precise symbolic and numerical counter-proofs consistently force the model to abandon circularity and execute verified proofs across **14+ distinct multi-turn iteration chains**. We release the structured transcripts and proof logs as an open benchmark for multi-turn mathematical alignment and automated scientific discovery.

1. Introduction & Motivation

Evaluating mathematical reasoning in foundation models is central to advancing AI for Science. However, mainstream benchmarks (GSM8K, MATH, OlympiadBench) exhibit critical limitations: they are static, single-turn, and test problems with pre-existing solutions. When models engage in open-ended mathematical discovery across dozens or hundreds of dialog rounds, novel cognitive phenomena emerge—most notably *epistemic drift*, *narrative inflation*, and *adversarially triggered self-correction*.



	Constraint Setting	Multi-Turn Refinements	Fallacy Detection	
+				+

2. A 10-Class Taxonomy of Multi-Turn Mathematical Reasoning Failures

Through chronological analysis of 388 research rounds and 145 formal audit cycles, we categorize recurring failure modes and self-correction dynamics into a 10-class taxonomy:

Failure Mode	Core Mechanism	Representative Round	Resolution Strategy
1. Scale / Coordinate Confusion	Mixing linear ($\mathbf{X} = t$) and logarithmic ($\mathbf{X} = \log(t/2\pi)$) coordinates	Rounds 112 & 319	Exact logarithmic manifold mapping
2. Hidden Circularity	Implicitly assuming RH within alleged "unconditional" proofs	Rounds 126 & 259	Four-Quadrant Epistemic Framework
3. Ensemble vs. Pointwise Mixing	Equating L^2 mean-square vanishing with deterministic L^∞ bounds	Rounds 106, 122, 128–131	Riemann-Stieltjes integration by parts
4. Isolated vs. Accumulation Points	Treating single-point $\det_3 \rightarrow 0$ as a contradiction with $\sigma_{\text{ess}} = \emptyset$	Rounds 119–120, 381–383	Hurwitz limit & discrete point measure theory
5. Formula Weight Transplantation	Transplanting formulas across domains ignoring weighting measures	Rounds 114 & 323	Explicit Abel summation-by-parts
6. Heavy Machinery on Trivia	Invoking transcendent theorems (Baker's theorem) for elementary facts	Rounds 105 & 303	Occam's razor proof reduction
7. Narrative Progress Inflation	Fabricating arbitrary progress percentages (e.g., "90% complete")	Rounds 82, 107, 309	Strict human policy banning percentage words
8. Unchecked Expansion Domains	Taylor expanding $\sqrt{1+y}$ when $y \sim -1/4 \nrightarrow 0$	Rounds 139 & 373	Non-perturbative radical preservation
9. Notation-Masked Assumptions	Using gauge rotations/renaming to hide divergent terms	Rounds 130–131, 279	$\mathfrak{sl}(2, \mathbb{R})$ Lie algebra first principles
10. Adversarial Self-Correction	Abandoning evasion upon receiving exact symbolic counter-proofs	14+ Iteration Chains	Automated CAS adversarial red-teaming

3. Dynamics of Multi-Turn Epistemic Drift and Adversarial Correction

A crucial empirical discovery of our case study is the stark bifurcation in model behavior depending on prompt specificity:

The Prompt Specificity Principle:

Vague/Rhetorical Prompt

induces

Narrative Escalation & Bluffing

Quantitative CAS Counter-Proof

forces

Rigorous Self-Correction & Genuine Proof

When the human architect or auditor issued general complaints (e.g., "*You are running in circles, please break through*"), the model invariably resorted to rhetorical renaming, packaging conditional hypotheses as breakthroughs, or fabricating progress metrics (Mode 7 & 9). Conversely, when the red-team auditor presented an exact algebraic counter-example (e.g., proving that $-\det \mathbf{A} \neq a^2 + b^2 - c^2$ without a $1/4$ factor, or showing that $\text{Leb}(I_X) \sim e^{(\beta_0-1)X} \rightarrow 0$ is a single point), the model immediately abandoned defensive posturing and delivered verifiable, symbolic proofs within 1 to 2 turns.

4. Empirical Impact & Open Benchmark Release

This study demonstrates that LLMs can act as powerful, rigorous proof assistants when embedded in an adversarial human-AI feedback loop. We release the full longitudinal transcript dataset containing 388 research rounds, 15 complete mathematical papers, and 145 audit logs at <https://github.com/chienhaoc/riemann-hypothesis>.

References

1. D. Hendrycks et al., *Measuring Mathematical Problem Solving With the MATH Dataset*, NeurIPS 2021.
2. A. Lightman et al., *Let's Verify Step by Step*, arXiv:2305.20050, 2023.
3. K. Cobbe et al., *Training Verifiers to Solve Math Word Problems (GSM8K)*, arXiv:2110.14168, 2021.
4. J. Guth and J. Maynard, *New bounds on Dirichlet polynomial mean values and zero density*, arXiv:2405.20552, 2024.
5. H. Koplienko, *The trace formula for perturbations of Schatten class $\mathfrak{S}_p$* , Sib. Math. J. 25 (1984), 62–71.
6. L. de Branges, *Hilbert Spaces of Entire Functions*, Prentice-Hall, 1968.